

# 1 # Comprehensive Assessment : Machine Learning

In [1]:

```
1 # Import necessary Libraries
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 from sklearn.model_selection import train_test_split, GridSearchCV
7 from sklearn.linear_model import LinearRegression
8 from sklearn.tree import DecisionTreeRegressor
9 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
10 from sklearn.svm import SVR
11 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

In [2]:

```
1 # Step 1: Load and Preprocess Data
2 #Loading and Preprocessing: The dataset is Loaded, and initial data exploration is performed. Missing values and categorical
3 url = 'https://drive.google.com/uc?id=1FHmYNLs9v0Enc-UExEMpit0FGsWvB2dP' # Direct Link
4 data = pd.read_csv(url)
```



In [3]:

```
1 # Explore the dataset
2 print(data.info())
3 print(data.describe())
4 print(data.isnull().sum())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 205 entries, 0 to 204
```

```
Data columns (total 26 columns):
```

#	Column	Non-Null Count	Dtype
0	car_ID	205 non-null	int64
1	symboling	205 non-null	int64
2	CarName	205 non-null	object
3	fueltype	205 non-null	object
4	aspiration	205 non-null	object
5	doornumber	205 non-null	object
6	carbody	205 non-null	object
7	drivewheel	205 non-null	object
8	enginelocation	205 non-null	object
9	wheelbase	205 non-null	float64
10	carlength	205 non-null	float64
11	carwidth	205 non-null	float64
12	carheight	205 non-null	float64
13	curbweight	205 non-null	int64
14	enginetype	205 non-null	object
15	cylindernumber	205 non-null	object
16	enginesize	205 non-null	int64
17	fuelsystem	205 non-null	object
18	boreratio	205 non-null	float64
19	stroke	205 non-null	float64
20	compressionratio	205 non-null	float64
21	horsepower	205 non-null	int64
22	peakrpm	205 non-null	int64
23	citympg	205 non-null	int64
24	highwaympg	205 non-null	int64
25	price	205 non-null	float64

```
dtypes: float64(8), int64(8), object(10)
```

```
memory usage: 41.8+ KB
```

```
None
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000

	curbweight	enginesize	boreratio	stroke	compressionratio
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	2555.565854	126.907317	3.329756	3.255415	10.142537
std	520.680204	41.642693	0.270844	0.313597	3.972040
min	1488.000000	61.000000	2.540000	2.070000	7.000000
25%	2145.000000	97.000000	3.150000	3.110000	8.600000
50%	2414.000000	120.000000	3.310000	3.290000	9.000000
75%	2935.000000	141.000000	3.580000	3.410000	9.400000
max	4066.000000	326.000000	3.940000	4.170000	23.000000

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

car_ID	0
symboling	0
CarName	0
fueltype	0
aspiration	0
doornumber	0
carbody	0
drivewheel	0
enginelocation	0
wheelbase	0
carlength	0
carwidth	0
carheight	0
curbweight	0
enginetype	0
cylindernumber	0
enginesize	0
fuelsystem	0
boreratio	0
stroke	0
compressionratio	0
horsepower	0
peakrpm	0

```
citympg      0
highwaympg   0
price        0
dtype: int64
```

```
In [6]: 1 print(data.columns)
```

```
Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
       'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
       'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
       'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
       'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
       'price'],
      dtype='object')
```

```
In [9]: 1 X = data.drop('price', axis=1)
        2 y = data['price']
        3
        4
```

```
In [10]: 1 # One-hot encoding for categorical variables
        2 X = pd.get_dummies(X, drop_first=True)
        3
```

```
In [11]: 1 #Model Implementation: Five regression models are instantiated and trained on the training set.
        2 # Step 2: Model Implementation
        3 # Split data into training and testing sets
        4 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
        5
        6
```

```
In [12]: 1 # Initialize models
        2 models = {
        3     'Linear Regression': LinearRegression(),
        4     'Decision Tree': DecisionTreeRegressor(),
        5     'Random Forest': RandomForestRegressor(),
        6     'Gradient Boosting': GradientBoostingRegressor(),
        7     'Support Vector Regression': SVR()
        8 }
        9
       10
```

```
In [13]: 1 # Train and evaluate models
        2 results = {}
        3 for name, model in models.items():
        4     model.fit(X_train, y_train)
        5     y_pred = model.predict(X_test)
        6     results[name] = {
        7         'R-squared': r2_score(y_test, y_pred),
        8         'MSE': mean_squared_error(y_test, y_pred),
        9         'MAE': mean_absolute_error(y_test, y_pred)
       10     }
       11
       12
```

```
In [14]: 1 # Step 3: Model Evaluation
        2 #Model Evaluation: The performance of each model is evaluated using R-squared, MSE, and MAE, and results are summarized.
        3 results_df = pd.DataFrame(results).T
        4 print(results_df)
        5
        6
```

	R-squared	MSE	MAE
Linear Regression	-1.239056	1.767601e+08	7280.667793
Decision Tree	0.887221	8.903252e+06	1881.918707
Random Forest	0.956571	3.428491e+06	1298.078000
Gradient Boosting	0.933773	5.228255e+06	1646.783382
Support Vector Regression	-0.101989	8.699543e+07	5707.167500

```
In [15]: 1 # Identify the best performing model
        2 best_model = results_df.loc[results_df['R-squared'].idxmax()]
        3 print(f"Best Model: {best_model.name}, R-squared: {best_model['R-squared']}")
        4
        5
```

Best Model: Random Forest, R-squared: 0.9565706087055967

```
In [16]: 1 # Step 4: Feature Importance Analysis
2 Feature Importance Analysis: Feature importances from the Random Forest model are displayed to understand the impact of each
3
4 # Feature importance for tree-based models
5 importances = models['Random Forest'].feature_importances_
6 feature_importance_df = pd.DataFrame(importances, index=X.columns, columns=['Importance']).sort_values('Importance', ascending=True)
7 print(feature_importance_df)
8
9
```

	Importance
enginesize	0.568230
curbweight	0.248291
highwaympg	0.053124
horsepower	0.042431
car_ID	0.020533
...	...
CarName_toyota corolla 1600 (sw)	0.000000
CarName_volkswagen rabbit	0.000000
CarName_volkswagen rabbit	0.000000
CarName_nissan nv200	0.000000
CarName_nissan note	0.000000

[190 rows x 1 columns]

```
In [17]: 1 # Step 5: Hyperparameter Tuning
2 Hyperparameter Tuning: The Random Forest model is tuned using Grid Search, and the performance of the tuned model is evaluated
3 # Example for Random Forest Regressor
4 param_grid = {
5     'n_estimators': [100, 200],
6     'max_depth': [None, 10, 20],
7     'min_samples_split': [2, 5, 10]
8 }
9 grid_search = GridSearchCV(RandomForestRegressor(), param_grid, cv=5)
10 grid_search.fit(X_train, y_train)
11
12
```

```
Out[17]:  GridSearchCV
  estimator: RandomForestRegressor
    RandomForestRegressor
```

```
In [18]: 1 # Evaluate the tuned model
2 best_rf_model = grid_search.best_estimator_
3 y_pred_tuned = best_rf_model.predict(X_test)
4 print(f"Tuned Random Forest R-squared: {r2_score(y_test, y_pred_tuned)}")
```

Tuned Random Forest R-squared: 0.9570222346614141